

6.7900 Fall 2024: Lecture Notes 1

Revision: 9/19/24 2:48PM

There was a lot of logistical information which you can find at `gradml.mit.edu`.

1 What is machine learning?

Broadly, machine learning is a set of methods for *using data to become better at a task*.

Problem formalization

- Input space : \mathcal{X} is often (but not necessarily) a vector space, often \mathbb{R}^d
- Output space : \mathcal{Y} is a fixed finite discrete set, for *classification* problems; it could be a continuous set, such as \mathbb{R}^d in a *regression* problem
- Training set : $\{(x^{(n)}, y^{(n)})\}_{n=1}^N$: set of N pairs of values, where $x^{(n)} \in \mathcal{X}$ and $y^{(n)} \in \mathcal{Y}$
- Decision rule : a function $h : \mathcal{X} \rightarrow \mathcal{Y}$.
- Loss function : $L : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$. This function is potentially *asymmetric*, so that $L(a, g)$ is the loss (badness) of predicting or guessing g when the actual label is a .

Called h for “hypothesis.”

Example problem: spam detection Our *data* is a set of pairs, of an email message and a binary value describing whether that message is spam. Our *task* is to predict whether future email messages are or are not spam. We represent each email message as a vector of *features*, $x^{(n)}$, for example with each $x_i^{(n)}$ having value 1 if some word i occurs in the n th message and 0 otherwise. The associated *label* $y^{(n)} \in \{0, 1\}$ indicates not-spam or spam.

What should our loss function be? Maybe something like:

$$L(a, g) = \begin{cases} 0 & \text{if } a = g \\ 1 & \text{if } a = 1 \text{ and } g = 0 \\ 15 & \text{if } a = 0 \text{ and } g = 1 \end{cases}$$

This encodes the idea that it's worse to say a message is spam when it is not (because we might miss an important message), than to say it is not spam when it is (because it is a mild annoyance).

Supervised learning problem: Given training set $\{(x^{(n)}, y^{(n)})\}_{n=1}^N$, find a decision rule $h : \mathcal{X} \rightarrow \mathcal{Y}$ that *performs well*.

To “perform well” is to make predictions on as-yet-unseen inputs that have low loss. That is, that $L(y^{(N+1)}, h(x^{(N+1)}))$ is low. But we don't know what these inputs will be! Probability to the rescue!

Probability is the language of uncertainty.

Assume training data and future data are drawn from some distribution $p(x, y)$. Now we can define *risk* of h to be the expected loss of using h to make a prediction on a new example:

$$\text{risk}(h) = \mathbb{E}[L(Y, h(X))]$$

Exercise: Show that, if $\{(X^{(N+m)}, Y^{(N+m)})\}_{m=1}^M$ are independent and identically distributed, the average expected loss of using h to make predictions on the next M points is equal to $\text{risk}(h)$.

Optimal h for classification We'd love to pick h to minimize risk $\mathbb{E}[L(Y, h(X))]$. If we know $p(X, Y)$, then we can!

Proposition: Consider K -class classification, in which $\mathcal{Y} = \{1, \dots, K\}$, and suppose X is a discrete random variable, Y is a random variable with domain \mathcal{Y} , and we know $p(x, y)$. Then the rule

$$h(x) = \arg \min_k \sum_{j=1}^K L(j, k) p(y = j \mid x)$$

minimizes $\mathbb{E}[L(Y, h(X))]$.

Proof. We start by expanding and rearranging terms in the risk definition:

$$\begin{aligned}
 \mathbb{E}[L(Y, h(X))] &= \sum_{x, y \in \mathcal{X} \times \mathcal{Y}} L(y, h(x)) p(x, y) \\
 &= \sum_{j=1}^K \sum_{x \in \mathcal{X}} L(j, h(x)) p(x, y = j) \\
 &= \sum_{j=1}^K \sum_{x \in \mathcal{X}} L(j, h(x)) p(y = j \mid x) p(x) \\
 &= \sum_{x \in \mathcal{X}} p(x) \sum_{j=1}^K L(j, h(x)) p(y = j \mid x)
 \end{aligned}$$

Now observe that this sum over x can be minimized by independently minimizing the j for each x , and that $h(x) = \arg \min_k \sum_{j=1}^K L(j, k) p(y = j \mid x)$ is exactly the h that does so. \square

Exercise: State and prove a similar result for continuous \mathcal{X} .

Two-class decision rule So, if we have two classes 0 and 1, and known $p(x, y)$, what is the form of the optimal h ? (Remember that the loss function is $L(a, g)$ where a is the true value of Y and g is the “guess” we are proposing to make.)

$$h(x) = \begin{cases} 0 & \text{if } \mathbb{E}[L(Y, 0)] < \mathbb{E}[L(Y, 1)] \\ 1 & \text{otherwise} \end{cases}$$

which can be rewritten as

$$\begin{aligned}
 h(x) &= \begin{cases} 0 & \text{if } L(0, 0)p(X = x, Y = 0) + L(1, 0)p(X = x, Y = 1) < L(0, 1)p(X = x, Y = 0) + L(1, 1)p(X = x, Y = 1) \\ 1 & \text{otherwise} \end{cases} \\
 &= \begin{cases} 0 & \text{if } L(0, 0)p(Y = 0 \mid X = x) + L(1, 0)p(Y = 1 \mid X = x) < L(0, 1)p(Y = 0 \mid X = x) + L(1, 1)p(Y = 1 \mid X = x) \\ 1 & \text{otherwise} \end{cases} \\
 &= \begin{cases} 0 & \text{if } L(1, 0)p(Y = 1 \mid X = x) < L(0, 1)p(Y = 0 \mid X = x) \\ 1 & \text{otherwise} \end{cases}
 \end{aligned}$$

Exercise: Why was it okay to turn $p(X = x, Y = 1)$ into $p(Y = 1 \mid X = x)$ in the derivation above?

Exercise: What happened to $L(0, 0)$ and $L(1, 1)$? Where would they go in the expression above? Why was it okay to omit them?

Optimal h for regression Now let's consider a regression problem where $\mathcal{Y} = \mathbb{R}$. If we know $p(x, y)$ we can also determine the optimal (minimum risk) h !

Proposition: Assume $\mathcal{X} = \mathbb{R}^d$ and $\mathcal{Y} = \mathbb{R}$, and we know $p(x, y)$. If we are using *squared loss* $L(a, g) = (a - g)^2$, then the decision rule

$$h(x) = \mathbb{E}[Y \mid X = x]$$

minimizes $\mathbb{E}[L(Y, h(X))]$.

Exercise: Complete the proof! It may be useful to make this (common, but counter-intuitive) move at some point:

$$y - h(x) = y - \mathbb{E}[Y \mid X = x] + \mathbb{E}[Y \mid X = x] - h(x)$$

Pretty cool! It says that the best prediction, under squared loss, is the conditional mean of Y given $X = x$.

Major obstacle We don't know $p(x, y)$!

But we do have training data! So, let's use it in the process of picking h . It can only be helpful if it is related, somehow, to the data that we are going to have to make predictions about in the future. A typical, strong assumption is that *all data*, including training data and future data, are independent and identically distributed.

Given this relationship between past and future data, we might think about approximating the *risk* by the *empirical risk* (average loss on the training data), observing that

$$\mathbb{E}[L(Y, h(X))] \approx \frac{1}{N} \sum_{n=1}^N L(y^{(n)}, h(x^{(n)}))$$

Exercise: Will this be the key to solving all our problems in machine learning?