

# 6.7900 Fall 2024: Lecture Notes 1

There was a lot of logistical information which you can find at `gradml.mit.edu`.

## 1 What is machine learning?

Broadly, machine learning is a set of methods for *using data to become better at a task*.

### Problem formalization

- Input space :  $\mathcal{X}$  is often (but not necessarily) a vector space, often  $\mathbb{R}^d$
- Output space :  $\mathcal{Y}$  is a fixed finite discrete set, for *classification* problems; it could be a continuous set, such as  $\mathbb{R}^d$  in a *regression* problem
- Training set :  $\{(x^{(n)}, y^{(n)})\}_{n=1}^N$  : set of  $N$  pairs of values, where  $x^{(n)} \in \mathcal{X}$  and  $y^{(n)} \in \mathcal{Y}$
- Decision rule : a function  $h : \mathcal{X} \rightarrow \mathcal{Y}$ .
- Loss function :  $L : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$ . This function is potentially *asymmetric*, so that  $L(a, g)$  is the loss (badness) of predicting or guessing  $g$  when the actual label is  $a$ .

Called  $h$  for “hypothesis.”

**Example problem: spam detection** Our *data* is a set of pairs, of an email message and a binary value describing whether that message is spam. Our *task* is to predict whether future email messages are or are not spam. We represent each email message as a vector of *features*,  $x^{(n)}$ , for example with each  $x_i^{(n)}$  having value 1 if some word  $i$  occurs in the  $n$ th message and 0 otherwise. The associated *label*  $y^{(n)} \in \{0, 1\}$  indicates not-spam or spam.

What should our loss function be? Maybe something like:

$$L(a, g) = \begin{cases} 0 & \text{if } a = g \\ 1 & \text{if } a = 1 \text{ and } g = 0 \\ 15 & \text{if } a = 0 \text{ and } g = 1 \end{cases}$$

This encodes the idea that it's worse to say a message is spam when it is not (because we might miss an important message), than to say it is not spam when it is (because it is a mild annoyance).

**Supervised learning problem:** Given training set  $\{(x^{(n)}, y^{(n)})\}_{n=1}^N$ , find a decision rule  $h : \mathcal{X} \rightarrow \mathcal{Y}$  that *performs well*.

To “perform well” is to make predictions on as-yet-unseen inputs that have low loss. That is, that  $L(y^{(N+1)}, h(x^{(N+1)}))$  is low. But we don't know what these inputs will be! Probability to the rescue!

Probability is the language of uncertainty.

Assume training data and future data are drawn from some distribution  $p(x, y)$ . Now we can define *risk* of  $h$  to be the expected loss of using  $h$  to make a prediction on a new example:

$$\text{risk}(h) = \mathbb{E}[L(Y, h(X))]$$

**Exercise:** Show that, if  $\{(X^{(N+m)}, Y^{(N+m)})\}_{m=1}^M$  are independent and identically distributed, the average expected loss of using  $h$  to make predictions on the next  $M$  points is equal to  $\text{risk}(h)$ .

**Optimal  $h$  for classification** We'd love to pick  $h$  to minimize risk  $\mathbb{E}[L(Y, h(X))]$ . If we know  $p(X, Y)$ , then we can!

**Proposition:** Consider  $K$ -class classification, in which  $\mathcal{Y} = \{1, \dots, K\}$ , and suppose  $X$  is a discrete random variable,  $Y$  is a random variable with domain  $\mathcal{Y}$ , and we know  $p(x, y)$ . Then the rule

$$h(x) = \arg \min_k \sum_{j=1}^K L(j, k) p(y = j \mid x)$$

minimizes  $\mathbb{E}[L(Y, h(X))]$ .

*Proof.* We start by expanding and rearranging terms in the risk definition:

$$\begin{aligned}
 \mathbb{E}[L(Y, h(X))] &= \sum_{x, y \in \mathcal{X} \times \mathcal{Y}} L(y, h(x)) p(x, y) \\
 &= \sum_{j=1}^K \sum_{x \in \mathcal{X}} L(j, h(x)) p(x, y = j) \\
 &= \sum_{j=1}^K \sum_{x \in \mathcal{X}} L(j, h(x)) p(y = j \mid x) p(x) \\
 &= \sum_{x \in \mathcal{X}} p(x) \sum_{j=1}^K L(j, h(x)) p(y = j \mid x)
 \end{aligned}$$

Now observe that this sum over  $x$  can be minimized by independently minimizing the  $j$  for each  $x$ , and that  $h(x) = \arg \min_k \sum_{j=1}^K L(j, k) p(y = j \mid x)$  is exactly the  $h$  that does so.  $\square$

**Exercise:** State and prove a similar result for continuous  $\mathcal{X}$ .

**Two-class decision rule** So, if we have two classes 0 and 1, and known  $p(x, y)$ , what is the form of the optimal  $h$ ?

$$h(x) = \begin{cases} 0 & \text{if } \mathbb{E}[L(0, h(X))] < \mathbb{E}[L(1, h(X))] \\ 1 & \text{otherwise} \end{cases}$$

which can be rewritten as

$$h(x) = \begin{cases} 0 & \text{if } L(1, 0)p(y = 1 \mid x) < L(0, 1)p(y = 0 \mid x) \\ 1 & \text{otherwise} \end{cases}$$

**Exercise:** What happened to  $L(0, 0)$  and  $L(1, 1)$ ? Where would they go in the expression above? Why was it okay to omit them?

**Optimal  $h$  for regression** Now let's consider a regression problem where  $\mathcal{Y} = \mathbb{R}$ . If we know  $p(x, y)$  we can also determine the optimal (minimum risk)  $h$ !

**Proposition:** Assume  $\mathcal{X} = \mathbb{R}^d$  and  $\mathcal{Y} = \mathbb{R}$ , and we know  $p(x, y)$ . If we are using *squared loss*  $L(a, g) = (a - g)^2$ , then the decision rule

$$h(x) = \mathbb{E}[Y \mid X = x]$$

minimizes  $\mathbb{E}[L(Y, h(X))]$ .

**Exercise:** Complete the proof! It may be useful to make this (common, but counter-intuitive) move at some point:

$$y - h(x) = y - \mathbb{E}[Y \mid X = x] + \mathbb{E}[Y \mid X = x] - h(x)$$

Pretty cool! It says that the best prediction, under squared loss, is the conditional mean of  $Y$  given  $X = x$ .

**Major obstacle** We don't know  $p(x, y)$ !

But we do have training data! So, let's use it in the process of picking  $h$ . It can only be helpful if it is related, somehow, to the data that we are going to have to make predictions about in the future. A typical, strong assumption is that *all data*, including training data and future data, are independent and identically distributed.

Given this relationship between past and future data, we might think about approximating the *risk* by the *empirical risk* (average loss on the training data), observing that

$$\mathbb{E}[L(Y, h(X))] \approx \frac{1}{N} \sum_{n=1}^N L(y^{(n)}, h(x^{(n)}))$$

**Exercise:** Will this be the key to solving all our problems in machine learning?